

Cancer Classification

```
library(caret)

## Loading required package: lattice
## Loading required package: ggplot2
library(corrplot)

## corrplot 0.84 loaded
library(ggplot2)
library(dplyr)

##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
library(keep)
library(glmnet)

## Loading required package: Matrix
## Loading required package: foreach
## Loaded glmnet 2.0-16
library(NeuralNetTools)
library(randomForest)

## randomForest 4.6-14
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
## The following object is masked from 'package:dplyr':
##
##   combine
## The following object is masked from 'package:ggplot2':
##
##   margin
library(tidyr)

## Warning: package 'tidyr' was built under R version 3.5.2
##
## Attaching package: 'tidyr'
## The following object is masked from 'package:Matrix':
##
```

```
##      expand
set.seed(1101)
```

read data

The breast cancer data consists of 30 features , they are computed from a digitized image of a fine needle aspirate (FNA) of a breast mass. They describe characteristics of the cell nuclei present in the image. n the 3-dimensional space is that described in: [K. P. Bennett and O. L. Mangasarian: “Robust Linear Programming Discrimination of Two Linearly Inseparable Sets”, Optimization Methods and Software 1, 1992, 23-34].

This database is also available through the UW CS ftp server: ftp ftp.cs.wisc.edu cd math-prog/cpo-dataset/machine-learn/WDBC/

Also can be found on UCI Machine Learning Repository: <https://archive.ics.uci.edu/ml/datasets/Breast+Cancer+Wisconsin+%28Diagnostic%29>

The target variable is diagnosis, tumor being malignant or benign. These 30 features are measures of the tumor such as radius, size, perimeter etc

```
bcancer <- read.csv("data.csv")
table(bcancer$diagnosis)
```

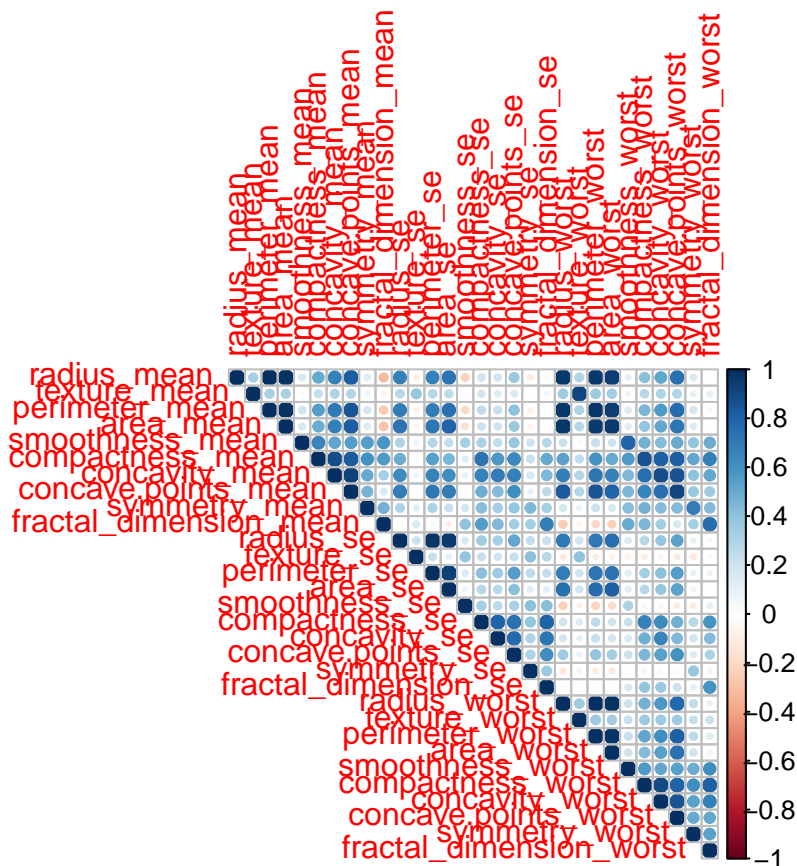
```
##
##      B      M
## 357 212
```

The objective of the analysis is to predict the the diagnosis of each patient id using these 30 features. I will use a classification model to identify the diagnosis.

DATA EXPLORATION

There are no missing values in the data and the distribution of the target variable is 63% of benign cancer and 37% of malignant cancer cells.

```
corMatrix <- cor(bcancer[,3:32])
corrplot(corMatrix , tl.cex = 1, addrect = 8 , type = "upper")
```



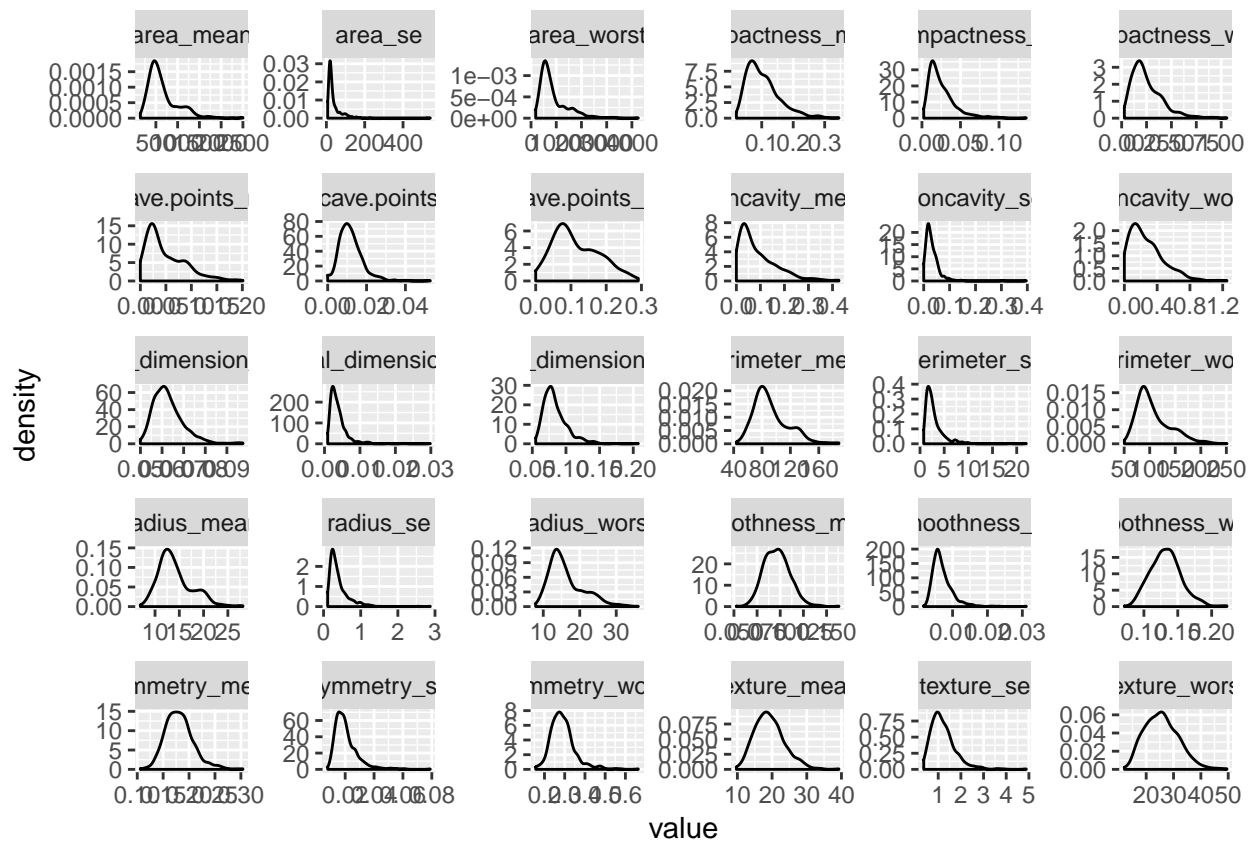
```
summary(bcancer[,3:32])
```

```
##      radius_mean      texture_mean      perimeter_mean      area_mean
##  Min.   : 6.981      Min.   : 9.71      Min.   : 43.79      Min.   : 143.5
##  1st Qu.:11.700      1st Qu.:16.17      1st Qu.: 75.17      1st Qu.: 420.3
##  Median :13.370      Median :18.84      Median : 86.24      Median : 551.1
##  Mean   :14.127      Mean   :19.29      Mean   : 91.97      Mean   : 654.9
##  3rd Qu.:15.780      3rd Qu.:21.80      3rd Qu.:104.10     3rd Qu.: 782.7
##  Max.   :28.110      Max.   :39.28      Max.   :188.50     Max.   :2501.0
##  smoothness_mean      compactness_mean      concavity_mean      concave.points_mean
##  Min.   :0.05263      Min.   :0.01938      Min.   :0.00000      Min.   :0.00000
##  1st Qu.:0.08637      1st Qu.:0.06492      1st Qu.:0.02956      1st Qu.:0.02031
##  Median :0.09587      Median :0.09263      Median :0.06154      Median :0.03350
##  Mean   :0.09636      Mean   :0.10434      Mean   :0.08880      Mean   :0.04892
##  3rd Qu.:0.10530      3rd Qu.:0.13040      3rd Qu.:0.13070      3rd Qu.:0.07400
##  Max.   :0.16340      Max.   :0.34540      Max.   :0.42680      Max.   :0.20120
##  symmetry_mean      fractal_dimension_mean      radius_se      texture_se
##  Min.   :0.1060      Min.   :0.04996      Min.   :0.1115      Min.   :0.3602
##  1st Qu.:0.1619      1st Qu.:0.05770      1st Qu.:0.2324      1st Qu.:0.8339
##  Median :0.1792      Median :0.06154      Median :0.3242      Median :1.1080
##  Mean   :0.1812      Mean   :0.06280      Mean   :0.4052      Mean   :1.2169
##  3rd Qu.:0.1957      3rd Qu.:0.06612      3rd Qu.:0.4789      3rd Qu.:1.4740
##  Max.   :0.3040      Max.   :0.09744      Max.   :2.8730      Max.   :4.8850
##  perimeter_se      area_se      smoothness_se      compactness_se
##  Min.   : 0.757      Min.   : 6.802      Min.   :0.001713      Min.   :0.002252
##  1st Qu.: 1.606      1st Qu.: 17.850      1st Qu.:0.005169      1st Qu.:0.013080
```

```
## Median : 2.287      Median : 24.530      Median :0.006380      Median :0.020450
## Mean   : 2.866      Mean   : 40.337      Mean   :0.007041      Mean   :0.025478
## 3rd Qu.: 3.357      3rd Qu.: 45.190      3rd Qu.:0.008146      3rd Qu.:0.032450
## Max.   :21.980      Max.   :542.200      Max.   :0.031130      Max.   :0.135400
## concavity_se      concave.points_se      symmetry_se
## Min.   :0.00000      Min.   :0.000000      Min.   :0.007882
## 1st Qu.:0.01509      1st Qu.:0.007638      1st Qu.:0.015160
## Median :0.02589      Median :0.010930      Median :0.018730
## Mean   :0.03189      Mean   :0.011796      Mean   :0.020542
## 3rd Qu.:0.04205      3rd Qu.:0.014710      3rd Qu.:0.023480
## Max.   :0.39600      Max.   :0.052790      Max.   :0.078950
## fractal_dimension_se radius_worst      texture_worst      perimeter_worst
## Min.   :0.0008948      Min.   : 7.93      Min.   :12.02      Min.   : 50.41
## 1st Qu.:0.0022480      1st Qu.:13.01      1st Qu.:21.08      1st Qu.: 84.11
## Median :0.0031870      Median :14.97      Median :25.41      Median : 97.66
## Mean   :0.0037949      Mean   :16.27      Mean   :25.68      Mean   :107.26
## 3rd Qu.:0.0045580      3rd Qu.:18.79      3rd Qu.:29.72      3rd Qu.:125.40
## Max.   :0.0298400      Max.   :36.04      Max.   :49.54      Max.   :251.20
## area_worst      smoothness_worst      compactness_worst      concavity_worst
## Min.   : 185.2      Min.   :0.07117      Min.   :0.02729      Min.   :0.0000
## 1st Qu.: 515.3      1st Qu.:0.11660      1st Qu.:0.14720      1st Qu.:0.1145
## Median : 686.5      Median :0.13130      Median :0.21190      Median :0.2267
## Mean   : 880.6      Mean   :0.13237      Mean   :0.25427      Mean   :0.2722
## 3rd Qu.:1084.0      3rd Qu.:0.14600      3rd Qu.:0.33910      3rd Qu.:0.3829
## Max.   :4254.0      Max.   :0.22260      Max.   :1.05800      Max.   :1.2520
## concave.points_worst symmetry_worst      fractal_dimension_worst
## Min.   :0.00000      Min.   :0.1565      Min.   :0.05504
## 1st Qu.:0.06493      1st Qu.:0.2504      1st Qu.:0.07146
## Median :0.09993      Median :0.2822      Median :0.08004
## Mean   :0.11461      Mean   :0.2901      Mean   :0.08395
## 3rd Qu.:0.16140      3rd Qu.:0.3179      3rd Qu.:0.09208
## Max.   :0.29100      Max.   :0.6638      Max.   :0.20750
```

Many features in the data are highly correlated. For e.g. radius mean and radius worst. This could cause multicollinearity in our models. Looking at the univariate plot of the data, we can see that the variables are very skewed. Most variables are right skewed and have varying range and scales. The variance in some of the area_se variable is quite high.

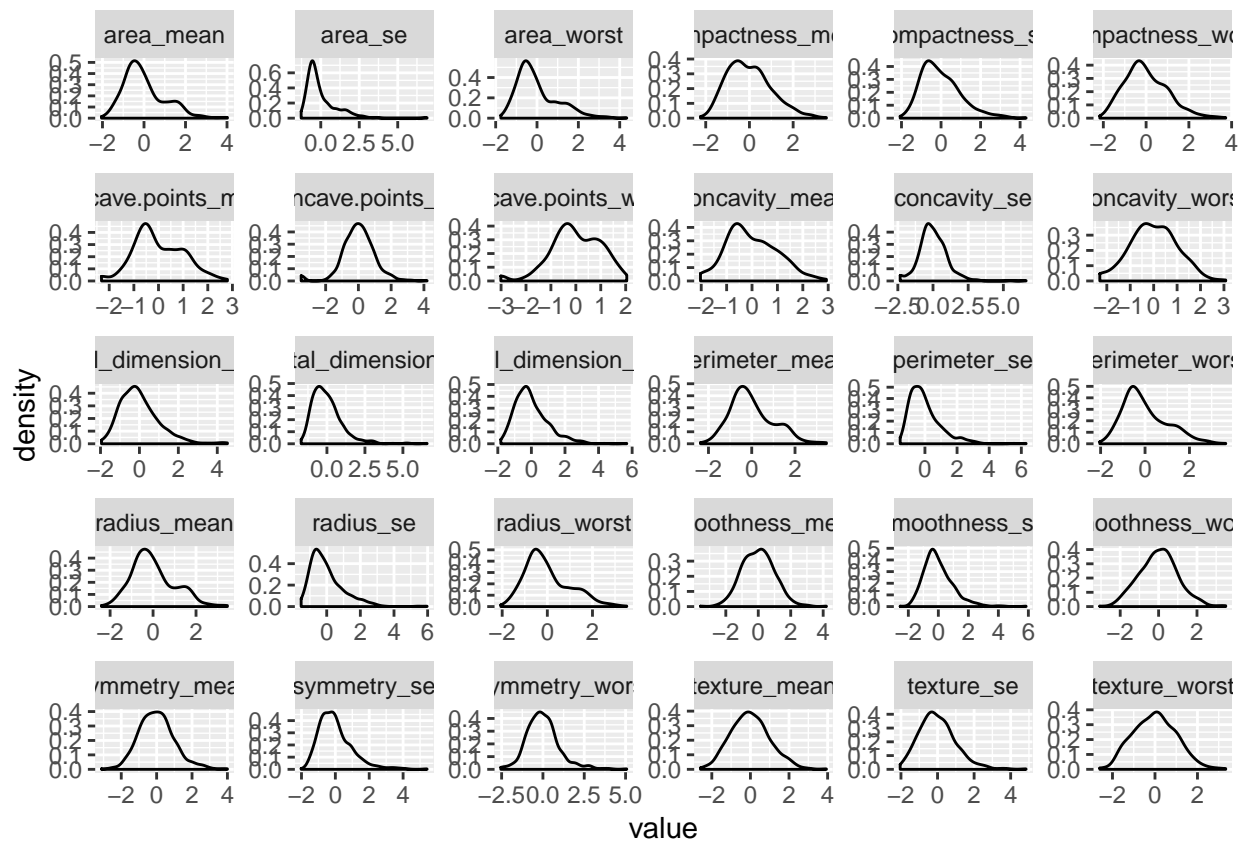
```
bcancer[,3:32] %>%
  #keep(is.numeric) %>%          # Keep only numeric columns
  gather() %>%                  # Convert to key-value pairs
  ggplot(aes(value)) +          # Plot the values
  facet_wrap(~ key, scales = "free") + # In separate panels
  geom_density()
```



As these distributions are skewed and vary largely on the scale I will transform the data. Using log transformation will cause errors as some values are 0 leading to undefined cases. Hence, I use a square root transformation and then rescale the data.

```
sdata <- bcancer
sdata[,3:32] <- sqrt(sdata[,3:32])
sdata[,3:32] <- scale(sdata[,3:32])

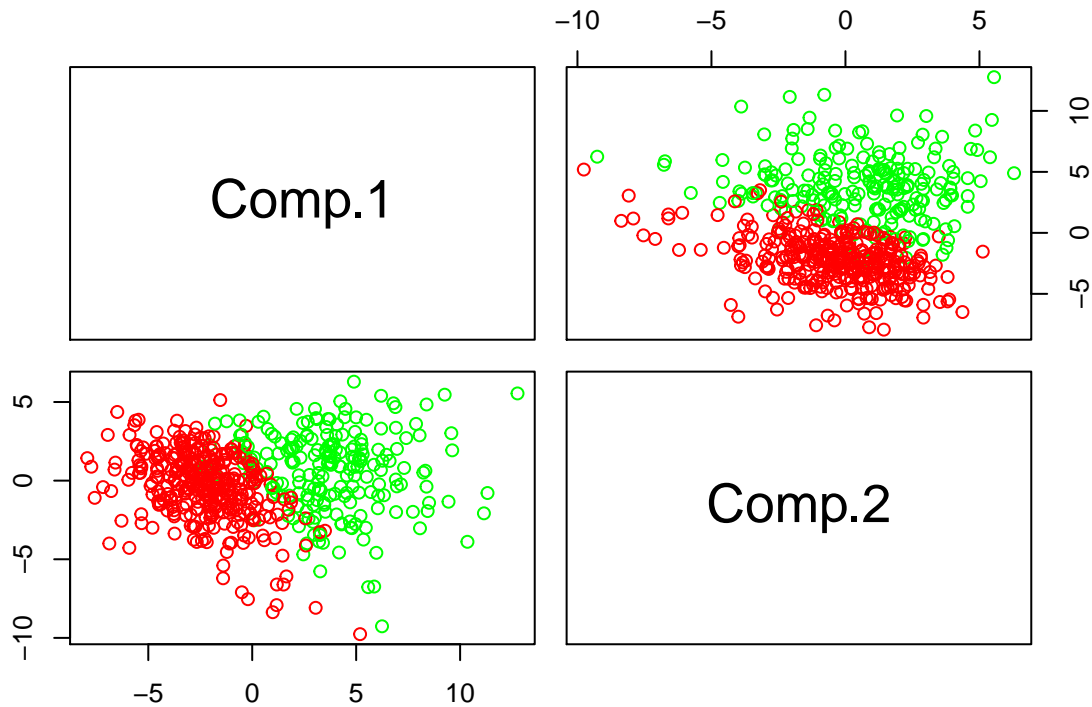
sdata[,3:32] %>%
  #keep(is.numeric) %>%           # Keep only numeric columns
  gather() %>%                   # Convert to key-value pairs
  ggplot(aes(value)) +           # Plot the values
  facet_wrap(~ key, scales = "free") + # In separate panels
  geom_density()
```



Now that the data is centered we could try different models on our data set. As there are many features in the data, to reduce the dimensionality of the data I will run a PCA on the scaled and transformed data.

```
pca <- princomp(sdata[,3:32])
pca_scores <- pca$scores
#pca$loadings

pairs(pca_scores[,1:2] , col = c("red" , "green")[sdata$diagnosis])
```



From the factor loading we can see that none of the components contribute heavily into the classification and only 7% variance is explained by the first 2 components. There are some overlap regions between the two classes which would be difficult to classify. Its hard to say from the princomp to decide which component will classify the data correctly, I chose to not go ahead with pca.

```
s <- sdata[,2:32]
inTrain <- createDataPartition(y=s$diagnosis, p=0.7, list=FALSE)
training <- s[inTrain,]
testing <- s[-inTrain,]
```

SVM - Linear

```
train_control <- trainControl(method="repeatedcv", number=10, repeats=20)
s <- sdata[,2:32]
svmLinear <- train(diagnosis~., data= training, trControl=train_control, method="svmLinear")

fit_svmLinear <- predict(svmLinear , newdata = testing)
cm_svmL <- confusionMatrix(fit_svmLinear , testing$diagnosis)
cm_svmL
```

Confusion Matrix and Statistics

```
##
##           Reference
## Prediction  B  M
##           B 107  7
##           M  0  56
##
##           Accuracy : 0.9588
##           95% CI : (0.917, 0.9833)
##           No Information Rate : 0.6294
##           P-Value [Acc > NIR] : < 2e-16
##
```

```
##           Kappa : 0.9097
## Mcnemar's Test P-Value : 0.02334
##
##           Sensitivity : 1.0000
##           Specificity : 0.8889
##           Pos Pred Value : 0.9386
##           Neg Pred Value : 1.0000
##           Prevalence : 0.6294
##           Detection Rate : 0.6294
##           Detection Prevalence : 0.6706
##           Balanced Accuracy : 0.9444
##
##           'Positive' Class : B
##
```

SVM - Radial

We will perform a 10 fold cross validation on the test data and do an out of sample testing for each model.

```
svmRadial <- train(diagnosis~., data= training, trControl=train_control, method="svmRadial")
```

```
fit_svmRadial <- predict(svmRadial , newdata = testing)
cm_svmR <- confusionMatrix(fit_svmRadial , testing$diagnosis)
cm_svmR
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  B  M
##           B 107  5
##           M   0 58
##
##           Accuracy : 0.9706
##           95% CI : (0.9327, 0.9904)
##           No Information Rate : 0.6294
##           P-Value [Acc > NIR] : < 2e-16
##
##           Kappa : 0.9359
## Mcnemar's Test P-Value : 0.07364
##
##           Sensitivity : 1.0000
##           Specificity : 0.9206
##           Pos Pred Value : 0.9554
##           Neg Pred Value : 1.0000
##           Prevalence : 0.6294
##           Detection Rate : 0.6294
##           Detection Prevalence : 0.6588
##           Balanced Accuracy : 0.9603
##
##           'Positive' Class : B
##
```

```
svmRadial$results$Accuracy
```

```
## [1] 0.9670545 0.9744583 0.9798365
```



```
svmRadial$results$AccuracySD
```

```
## [1] 0.02731376 0.02417758 0.02079046
```

SVM does a good job in predicting the classes with only 3 data points misclassified in the Radial SVM model. The insample accuracy of the model is also quite high with the standard deviation of 0.02. We can say the accuracy estimate of SVM is strong. In this problem, the false negatives are of high importance. It is crucial to detect the tumor so that the patients get treatments. Hence, we will try to reduce the false negatives, i.e those cases that are originally malignant but classified as benign.

```
knn <- train(diagnosis~., data= training, trControl=train_control, method="knn")
fit_knn <- predict(knn , newdata = testing)
CM_KNN <- confusionMatrix(fit_knn , testing$diagnosis)
CM_KNN
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction   B    M
##           B 105    5
##           M   2   58
##
##              Accuracy : 0.9588
##              95% CI : (0.917, 0.9833)
##      No Information Rate : 0.6294
##      P-Value [Acc > NIR] : <2e-16
##
##              Kappa : 0.9109
##  Mcnemar's Test P-Value : 0.4497
##
##              Sensitivity : 0.9813
##              Specificity : 0.9206
##              Pos Pred Value : 0.9545
##              Neg Pred Value : 0.9667
##              Prevalence : 0.6294
##              Detection Rate : 0.6176
##      Detection Prevalence : 0.6471
##              Balanced Accuracy : 0.9510
##
##              'Positive' Class : B
##
```

KNN also has high accuracy but again there are more false negatives , which is not desirable.

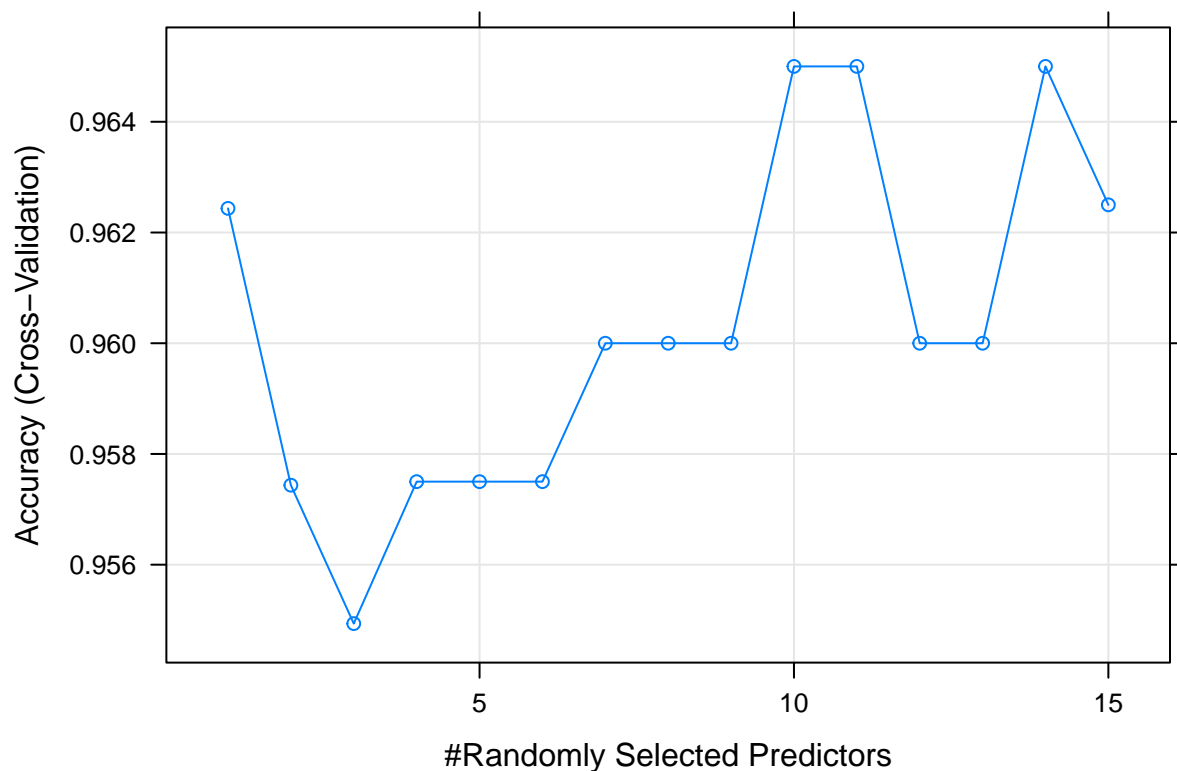
We will try thr random forest model with grid search approach.

```
tuneGrid <- expand.grid(.mtry = c(1:15) )
trControl <- trainControl(method = "cv", number = 10, search = "grid")
rf <- train(diagnosis~., data= training ,method = "rf", metric = "Accuracy", trControl =trControl,tuneGrid=tuneGrid)
pred_rf <-predict(rf, testing)
CM_RF <- confusionMatrix(pred_rf , testing$diagnosis)
CM_RF
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction   B    M
```

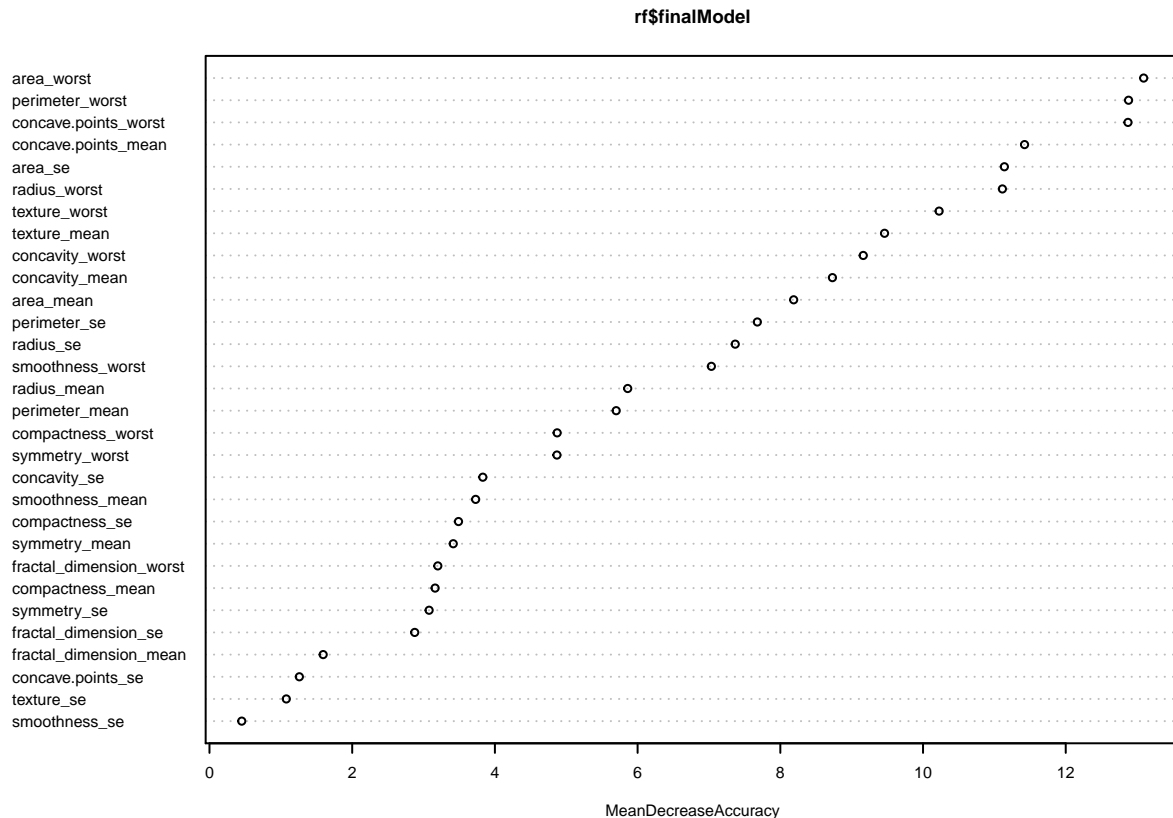
```
##          B 104   4
##          M   3  59
##
##          Accuracy : 0.9588
##          95% CI : (0.917, 0.9833)
##          No Information Rate : 0.6294
##          P-Value [Acc > NIR] : <2e-16
##
##          Kappa : 0.9114
##          Mcnemar's Test P-Value : 1
##
##          Sensitivity : 0.9720
##          Specificity : 0.9365
##          Pos Pred Value : 0.9630
##          Neg Pred Value : 0.9516
##          Prevalence : 0.6294
##          Detection Rate : 0.6118
##          Detection Prevalence : 0.6353
##          Balanced Accuracy : 0.9542
##
##          'Positive' Class : B
##
```

```
plot(rf)
```



The results of random forest are more promising with just 2 false negatives. There are still 2 malignant cases that are classified as benign. We also have increased number of false positive in this model. I tried out different mtry for grid search and we can see that the maximum accuracy is obtained when 7 features are randomly sampled for each evaluation which also close to $\sqrt{30}$

```
varImpPlot(rf$finalModel,type=1 ,cex=.5)
```



We also look at the variable importance to understand the contribution of each variable in the classification process. To try an improve our model I will next try a neural network.

```
set.seed(1101)
tuneGrid <- expand.grid(.size = c(1:6), .decay=c(0,2.5e-2,5e-2,7.5e-2,1e-1,1e-2) )
nnet <- capture.output(nn <- caret::train(diagnosis~., data= training, method = "nnet", metric = "Accuracy"))

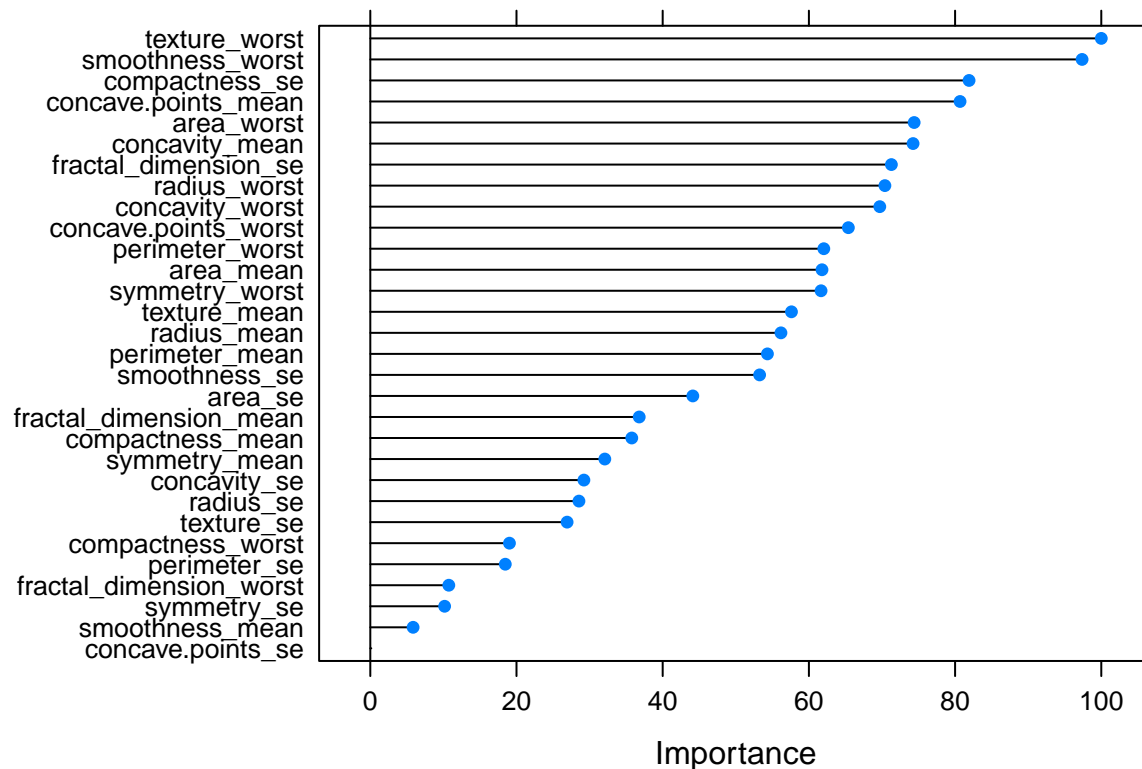
fit_nn <- predict(nn , testing)
t <- confusionMatrix(fit_nn , testing$diagnosis)
t
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction  B   M
##           B 106   4
##           M   1  59
##
##           Accuracy : 0.9706
##           95% CI : (0.9327, 0.9904)
##           No Information Rate : 0.6294
##           P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.9363
##           McNemar's Test P-Value : 0.3711
##
```

```
##          Sensitivity : 0.9907
##          Specificity : 0.9365
##          Pos Pred Value : 0.9636
##          Neg Pred Value : 0.9833
##          Prevalence : 0.6294
##          Detection Rate : 0.6235
##          Detection Prevalence : 0.6471
##          Balanced Accuracy : 0.9636
##
##          'Positive' Class : B
##
```

```
nnet_vatimpt <- varImp(nn)
plot(nnet_vatimpt)
```

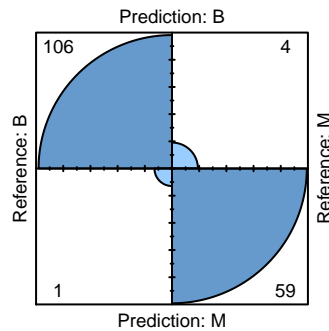


The neural network model also gives high accuracy and the missclassified rate is also low. Both the models random forest and Neural Network work well for this data. But I would like to go ahead with NNet model as it is less computational expensive compared to random forest.

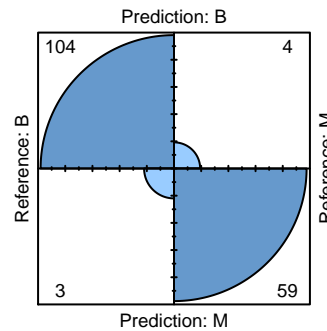
```
par(mfrow=c(2,2))

fourfoldplot(t$table ,conf.level = 0, margin = 1 , main = "Neural Net")
fourfoldplot(CM_RF$table ,conf.level = 0, margin = 1 , main = "Random Forest")
fourfoldplot(CM_KNN$table,conf.level = 0, margin = 1 , main = "KNN")
fourfoldplot(cm_svmR$table ,conf.level = 0, margin = 1 , main = "Linear SVM")
```

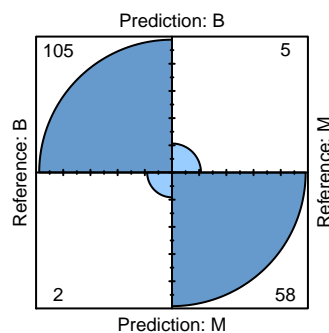
Neural Net



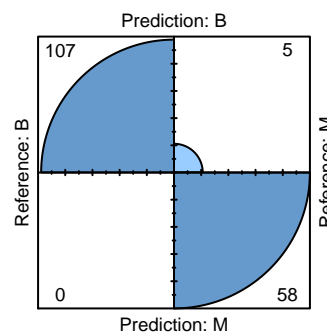
Random Forest



KNN



Linear SVM



Analysing the missclassified data to understand the error in our model. Through bivariate plots it is difficult to understand the error. Some expertise on the data and knowledge of the subject would probably help us understand the problem and build a finer model to detect the cancer.

```
misclassified <- testing[which(fit_nn != testing[,1]), ]
misclassified
```

```
##      diagnosis radius_mean texture_mean perimeter_mean area_mean
## 74           M -0.03606509 -0.81124238 -0.002097492 -0.1028500
## 136          M -0.34756777  0.77621175 -0.386399045 -0.3688263
## 216          M -0.01828376 -0.51937386  0.020676442 -0.1200526
## 264          M  0.48438093  0.07632073  0.399369535  0.4355538
## 364          B  0.72920285 -0.18375347  0.665137478  0.6600813
##      smoothness_mean compactness_mean concavity_mean concave.points_mean
## 74      0.34321220      0.5685772      0.08907512      0.2517173
## 136     -0.38574209     -0.9443617     -0.38190215     -0.4400026
## 216      0.47551577      0.9760872      0.35911332      0.3832229
## 264     -1.31440991     -0.9834196     -0.47216534     -0.3910957
## 364      0.07195321     -0.2893660     -0.19160176      0.1918038
##      symmetry_mean fractal_dimension_mean radius_se texture_se
## 74      -0.5288348      0.4396747 -0.4519290 -1.24250505
## 136     -0.8301097      -0.2883112 -0.6796914  0.41341433
## 216      1.0865208      0.9318799 -0.5710834  0.06242961
## 264     -0.9814880      -1.2353503 -0.7189902 -0.33678578
## 364     -1.1916831      -1.0022117 -0.1535941  0.51973648
##      perimeter_se area_se smoothness_se compactness_se concavity_se
## 74      -0.4568406 -0.39407448 -0.8500412 -0.1549129 -0.4492374
## 136     -0.8507038 -0.54233489  0.2635577 -0.8319450 -0.1635229
## 216     -0.4745039 -0.42129666 -0.3204383  0.6742002  0.4767663
```

```
## 264 -0.7860712 -0.44260427 -1.8147117 -1.1144742 -0.7071624
## 364 -0.1856965 -0.01308313 0.1759969 -0.3407554 -0.4104797
## concave.points_se symmetry_se fractal_dimension_se radius_worst
## 74 -0.2837356 -1.18118889 -0.1858153 0.13602251
## 136 -0.3340067 -0.48928137 -0.4569881 -0.32625021
## 216 0.5142158 -0.06704247 0.4863016 -0.04254602
## 264 -1.0875033 -1.56060319 -1.3042620 0.41857230
## 364 -0.0933569 -0.41892717 -0.8366042 0.46393764
## texture_worst perimeter_worst area_worst smoothness_worst
## 74 -0.77038345 0.16919269 0.007198285 0.4207874
## 136 1.23708790 -0.41966855 -0.339458820 0.4548136
## 216 0.26229439 -0.01534565 -0.124366960 0.6277119
## 264 0.98965394 0.33983417 0.354132653 -1.0636702
## 364 0.02225315 0.37885383 0.392226809 0.1057293
## compactness_worst concavity_worst concave.points_worst symmetry_worst
## 74 0.7683540 0.23384268 0.4739518 -0.4901866
## 136 -0.6292397 -0.05728276 -0.1483584 -0.0708075
## 216 1.2179557 0.97257653 0.8000677 1.2132903
## 264 -0.3916455 -0.01493833 -0.2678673 -0.3237052
## 364 -0.4962290 -0.33978416 -0.1804370 -0.8454749
## fractal_dimension_worst
## 74 1.1103462
## 136 -0.1430000
## 216 1.2626932
## 264 -0.9138259
## 364 -1.1508780
```

```
plot(testing$texture_worst, col=ifelse(rownames(testing) %in% rownames(misclassified[1,])
, 'red', c("green", "blue")[testing$diagnosis]), lower.panel = NULL)
legend(120, 3, legend = c("Malignant", "Benign", "Misscd as B"), col = c("green", "blue", "red"))
```

